

COMPACT REPRESENTATION OF SEPARABLE GRAPH

by

Xiang Zhang

Submitted in partial fulfillment of the requirements
for the degree of Master of Applied Computer Science

at

Dalhousie University
Halifax, Nova Scotia
August 2017

© Copyright by Xiang Zhang, 2017

Contents

Abstract	iii
Acknowledgements	iv
Chapter 1 Introduction	1
1.1 Related Works	3
1.2 Preliminaries	3
1.3 Outline	7
Chapter 2 Edge Separator Tree	8
Chapter 3 Index Structure	16
Chapter 4 Experiment	19
Chapter 5 Conclusion	21
Bibliography	22

Abstract

This is a test document.

Acknowledgements

My Supervisor, Dr. Meng He, deserves the most praise for his invaluable guidance, friendship and suggestions. I would also like to thank all of the friends and acquaintances I have made at Dalhousie University. They have made this journey something to remember. Without everyone's support and guidance this project would have not been a success.

Chapter 1

Introduction

Nowadays, many applications use graphs to show the relationship and represent connectivity between multiple objects. As the graphs inevitably grow very huge, the space issue becomes ever more important.

In the field of Computer Science, a succinct representation of a data structure is a compact representation whose space cost is close to the information-theoretic lower bound, and still supports efficient query operations. The concept of “succinct” was firstly introduced by Jacobson [2] to encode data structures such that bit vectors, (unlabeled) trees, and planar graphs.

In this project, we are interested in improving a compact representation mechanism for separable graphs presented by a previous work [3]. In which work the authors proposed an approach to representing separable graphs compactly. Their representation used $O(n)$ bits, meanwhile using constant time on degree or adjacency query, and neighbour listing for one vertex in constant time per neighbour (They took advantage of $O(\log n)$ -bit parallelism computation to access $O(n)$ consecutive bits in one operation). Graphs with good separators got good compression by using

their representation, and even graphs that are not strictly separable, their representation still works well because the separable components in those graphs can be compressed. In their paper, they provided detail description for compressing the graph by building two structures: the adjacency table and the root-find structure [3], which used vertex separators to encode the graph into a shadow adjacency table [3], as well as support constant time on query operations. But their experiment implemented the data structure by using edge separator instead of vertex separator, which means the shadow label and root-find structure were not needed. In this project, we implement the data structure by using edge separator, and conduct experiment on it.

We firstly follow their idea to implement the data structure by using edge separator. The data structure is building by recursively partitioning a graph into two subgraphs until only one vertex left. During the partition process, an edge separator tree is built. Then make use of the edge separator tree to renumber the vertexes. We store an adjacency list for each vertex, then concatenate the adjacency lists in the order of renumbered vertex to form an adjacency table. Each pair of vertexes in the adjacency is stored via encoding the difference d between the two vertexes, and all the differences are stored contiguously as a sequence of bits in memory. We show, by using adjacency table, $O(n)$ bits are sufficient to encode the separable graph.

Next, we implement several index structures to support degree and adjacency query in constant time, and neighbour listing in constant time per Neighbour. The index structures contains all the index structures involved in their experiment, as well

as two index structures implemented by us via applying succinct data structure on bit vector. In their paper, they encoded the degree of each vertex at the start of each adjacency list. However, we found that, the space of degree can be saved in some cases, which causes a trade-off between the time to support degree queries and the space to encode the graph.

1.1 Related Works

1.2 Preliminaries

Graph Separator. A family of graphs G is define separable if : it is closed under taking subgraphs, and satisfies the $f(\cdot)$ -separator theorem [8] if there is a constant $\alpha < 1$ and $\beta > 0$ such that each member graph in G with n vertexes has a separator set S of size $|S| < \beta f(n)$, that partition the graph into two parts A and B , with at most αn vertexes in each ‘[8]. A graph is separable if it belongs to a separable family of graphs.

In this project we focus on the graphs that satisfy the n^c -separator theorem for some constant $c < 1$. One class of graphs that reach this specifications is the planar graphs, which satisfies the theorem that $c = \frac{1}{2}$. Another example can be well-shaped meshes in \mathbb{R}^d , with separators of size $O(n^{1-1/d})$ [5].

Lipton *etal.* [4] prove that all classes of graph which satisfy $n/(\log n)^{1+\epsilon}$ -separator theorem have bounded density. The bounded density means that every n -vertex

member in a class if graphs has bounded $O(n)$ edges. So we can assert that the separable graphs have bounded density.

By making use of the definitions above, we define a class of graphs G satisfies a $f(n)$ -edge separator theorem if there are constant variables $\alpha < 1$ and $\beta > 0$ such that each member graph in G with n vertexes, has an edge separator with at most $\beta f(n)$ edges whose removal partitions the graph into two subgraphs with at most αn vertexes in each. The edge separator is not as common as the vertex separator, because a graph with an edge separator of size s also has a vertex separator of size s at most, but no similar bounds holds when it is conversed [3].

In this project, we will only consider the connected graph, which means all vertexes in the graph have nonzero degree, and we assume each step of partitioning always returns a edge separator set of size $O(n^c)$.

Queries. Our data structure supports three kinds of queries on separable graphs: adjacency queries, degree queries and neighbour listing. The adjacency query checks whether there is an edge between two vertices, the neighbour listing returns the neighbours of a given vertex, and the degree query reports the number of edges connected to a given vertex.

Bit Vector. The compact representation stores a separable graph as difference-encoded bit sequence in memory. To support the three kinds queries (adjacency

queries, degree queries and neighbour listing) in constant time, knowing both the number of encoded adjacency lists in the sequence up to an index position, and the start position of a particular adjacency is necessary. More formally, the compact representation requires data structures that support rank and select operations in constant time: Given a set $B[0..n]$ to represent a subset S of a universe $U = [0..n) = \{0, 1, \dots, n-1\}$, where $B[i] = 1$ iff $i \in S$,

- $rank_q(x) = \{k \in [0..x] : B[k] = q\}$
- $select_q(x) = \min\{k \in [0..n) : rank_q(k) = x\}$

In this project, we use two kinds of succinct data structure on bit vector : RRR [7] and SD vector [6], to build the index structure which supports queries on separable graph. Both structures satisfy information-theoretic lower bound, and still supports efficient rank and select query operations.

Adjacency Tables. An adjacency list representation for a graph associates each vertex in the graph with the collection of its neighboring vertices. For use as a data structure, the main alternative to the adjacency list is the adjacency matrix. For a sparse graph (most pairs of vertices are not connected by edges), in our case, a graph has *bounded mdensity*, an adjacency list representation is significantly more space-efficient than an adjacency matrix representation: the space cost of the adjacency list representation depends on the number of edges and vertices in the graph, while the adjacency matrix representation is stored as the square of the number of vertices. Generally, an adjacency list representation for an undirected graph requires $8|E|$ bytes

($2(32)|E|$ bits) of space, where $|E|$ is the number of edges in the graph. However, it is not space-efficient enough yet.

In our project, we assign each vertex in the graph an integer label. For each vertex, we store an adjacency list which contains its all neighbouring vertexes. If a vertex with label v has neighbours $v_1, v_2, v_3, \dots, v_n$ in ascending sorted order, then we encode the difference between each adjacency pair in the adjacency list as $|v_1 - v|, |v_2 - v_1|, |v_3 - v_2|, \dots, |v_n - v_{n-1}|$ contiguously, which forms a sequence of bits stored in memory. We use the gammacode [1] to encode the differences, which uses $2\lfloor \log n \rfloor + 1$ bits to encode a difference. The difference $v_1 - v$ may be negative in some cases, we here store a one-bit flag that value. To implement one of our index structures, the degree of each vertex needs to be encoded at start position of the corresponding adjacency, but in other circumstances the space of degree can be saved. By concatenating all the adjacency lists in the order of the vertex labels, we form an adjacency table. To access the adjacency list of a vertex, knowing the starting position of the adjacency is necessary. The easiest and fastest approach is to store an array of offset pointers of size $O(\log(n))$ for each. But if doing so, we would use $nO(\log(n))$ space, which exceeds our space bound. Alternatively, we implement other two index structures proposed by Blandfor [3], as well as two index structures based on succinct data structure on bit vector, which use less space cost down to $O(n)$ bits.

Lemma 1.2.1. *An adjacency table can support degree queries in $O(1)$ time, and neighbour listing in $O(d)$ time, where d is the degree of the vertex.*

Proof. To access an adjacency list of a particular vertex, we use our index structure to locate its starting position in constant time. Each adjacency list consists of a encoded degree number and a sequence of encoded differences, which use $O(\log(n))$ bits and $dO(\log(n))$ bits respectively. By taking advantages of the $O(\log n)$ -bits parallelism computation, we can decode any $O(\log(n))$ -bits value in constant time, like using lookup table. Hence we can decode the degree at the starting of each adjacency list in constant time. For neighbour listing, we decode each difference pair in constant time. So it takes constant time to make degree queries, and $O(d)$ time to make neighbourhood queries, where d is the degree of the vertex. \square

1.3 Outline

This report is organized as follows. Chapter 1 shows the motivation of this project, and provides some preliminaries used in this project. Chapter 2 and 3 explain the construction and usage of the edge separator tree and several index structures respectively. Chapter 4 describes the experiment setup and examines the performance of our data structure. Chapter 5 shows the conclusion as well as the future work of our project.

Chapter 2

Edge Separator Tree

To compactly represent the tree, the first thing we need to do is building an edge separator tree from the original tree. The tree-building process is based on recursively partition the tree into two parts using edge separator. One example that illustrates the function of edge separator is shown in Figure 2.1: assuming we have an 8-vertex graph G_1 with 9 edges, by removing the two edges (v_1, v_9) and (v_1, v_5) , we can partition G_1 into two parts with 4 vertexes in each. Hence the edge separator in this partitioning step is $\{(v_1, v_9), (v_1, v_5)\}$.

We can see that each time we make a partitioning on a graph, it splits its vertex set and parts of edges set, and removes some edges. if we perform this partitioning recursively until only one vertex left, then we have an edge separator tree from the graph, one edge separator tree for the G_1 in Figure 2.1 is shown in Figure 2.2. Each vertex in a graph will appear once in a leaf of the edge separator tree build from that graph. Each internal node is the set of edges used as edge separators to make partitioning in each step.

Note that we arbitrarily decide which side of partition will be the left child or the right child in this stage. The tree is also not perfect balanced in some cases (e.g.




Figure 2.1: An example of a partition process using edge separator

one child is a leaf and the other is an internal node), the side of the children is also decided arbitrarily. We will take advantages of this "freedom" to further reduce the space to encode the graph in later time. The vertex labels may be disordered by making in-order traversal on the edge separator tree (e.g. in Figure 2.2 the order of the vertexes after in-order traversal is $:v_1, v_4, v_2, v_3, v_5, v_6, v_7, v_8$), we will renumber these vertex by making a in-order traversal along the edge separator tree.

METIS Partition Library. To recursively partition the graph, we use METIS's API in our project. METIS is an open source library developed by Karypis lib in University of Minnesota. It supports set of serial programs for partitioning graphs and partitioning finite element meshes. The algorithms implemented in METIS are based on the multilevel recursive-bisection, multilevel k-way, and multi-constraint




Figure 2.2: An example of edge separator tree

partitioning schemes [2]. The METIS partitioning API returns one integer list with vertex index as key, and a flag 0 or 1 to indicate which side of partitioning the vertex belongs to. The returned value of METIS API does not provide information about edges, and the API takes two integer arrays (adjacency list of current graph and the index of vertex in adjacency list) as input parameters, hence before each calling of the METIS API, we need to generate an adjacency list and an index list for the current graph to be partitioned, and in order to collect the edge separators, we need both the adjacency list and the returned values after calling the API.

The Algorithm to build the separator tree is given in Algorithm 1. Here we use (V, E) to represent a graph for simplicity, but in fact the METIS API takes formalized adjacency list and the index list as input parameters. So each time to call the METIS

API, the current vertex set need to be renumbered temporally in contiguous ascending numbers, and the adjacency list needs to be initialized according to the temporally order. The METIS API also does not support to partition the graph with the vertex amount under 3, hence the very basic partitioning cases are defined by ourself.

Algorithm 1: Building the edge separator tree

```

1 buildTree ( $V, E$ );
2 if  $|V| = 1$  then
3   | return  $V$ ;
4 else
5   |  $(V_a, V_b) \leftarrow \text{METIS\_API}(V, E)$  ;
6   |  $E_{sep} \leftarrow \text{findEdgeSeparator}(V, E, V_a, V_b)$  ;
7   |  $E' = E - E_{sep}$  ;
8   |  $E_a \leftarrow \{(u, v) \in E' \mid u \in V_a \vee v \in V_b\}$  ;
9   |  $E_b = E' - E_a$  ;
10  |  $T_a \leftarrow \text{buildTree}(V_a, E_a)$  ;
11  |  $T_b \leftarrow \text{buildTree}(V_b, E_b)$  ;
12  | return separatorTree ( $T_a, E_{sep}, T_b$ ) ;
13 end

```

Building Adjacency Table. After getting the edge separator tree, we need to renumber the vertexes based on an in-order traversal along the tree. Each vertex appears in one leaf of the tree will receive a new label which increases by one when visiting a new leaf node. A mapping table is also generated after completing the traversal, which will be used to build an adjacency list for each renumbered vertex by using the new vertex labels. Then we can build an adjacency list for each vertex, the list contains all the neighbouring vertexes in sorted order. We use gamma code to encode the difference of each adjacency pair in the list (e.g. $v_1 - v$, $v_2 - v_1, \dots$, $v_{n-1} - v_{n-2}$), each difference d costs $2\lceil \log n \rceil + 1$ bits to be encoded. Note that the

value $v - v_1$ might be negative, so we store one flag bit for it. We also encode the number of the entries in an adjacency list at the starting position of the list, which will be used to support degree queries in constant time.

Lemma 2.0.1. *Any n -vertex member in a class of graphs which satisfies a n^c – edge separator theorem, can be encoded in $O(n)$ bits using an adjacency table.*

Proof. For each edge (u, v) in an edge separator in a graph with s vertexes, the difference between vertex u and vertex v is $O(s)$. We use gamma code to encode the difference which would use $O(\log(s))$ bits, so that edge will contribute $O(\log(s))$ bits to the adjacency lists of both vertex u and vertex v . Hence, we charge $O(\log(s))$ to every edge in a separator of a graph with s vertices. Recall that by using the n^c – edge separator theorem, each member with n vertexes in the class of the graph has a βn^c edge separator whose removal partitions the member into two parts with at most βn vertexes in each. Let define $S(n)$ to be an upper bound of required bits to encode a graph with n vertexes, and let $\alpha < a < 1 - \alpha$, then $S(n)$ satisfies the recurrence:

$$S(n) \leq S(an) + S(n - an) + O(n^c \log n)$$

This recurrence can be solved to $S(n) = O(n)$, and the bits to encode the degree of all the vertexes is bounded by $O(n)$, because there are $O(n)$ edges in total. So the total space is $O(n)$ bits. □

Child-flipping. When encoding the graph, We contiguously encode the difference

between each adjacency pair in the list by using gamma code. The gamma code use $2\lceil \log n \rceil + 1$ bits to encode a difference d . If we can make each difference "smaller", in other words, if we can make the labels of vertexes in an adjacency list closer to each other, then we can reduce the space to encode the adjacency list. We arbitrarily decide the side of partitioning during the construction of the separator tree (in our project, we decide the side of partitioning according to the output value of METIS API: '0' to the left child and '1' to right child), so there is a degree of freedom in the way to build the tree. To take advantage of this freedom, we can apply a heuristic optimization method called "child-flipping" [3] to our edge separator tree.

The child-flipping algorithm works when making traversal along a separator tree, it tracks the nodes containing vertex labels which appear before and after the vertex labels in the current node. Let define N_L to be the left children of current node's left ancestors and N_R to be the right children of the current node's right ancestors, N_1 and N_2 represents the current node's left child and right child, respectively, and E_{AB} indicates the number of edges between the vertexes in node A and node B. For each node, the vertexes in N_1 of the current node will have labels less than the vertexes' under the node, the vertexes in N_2 will have larger labels than all the vertexes under the node, and there is no intersection between the vertexes in N_L and N_R of the node. The child-flipping exams $E_{N_L N_1}$, $E_{N_2 N_R}$, $E_{N_L N_2}$ and $E_{N_1 N_R}$ of each node to ensure that $E_{N_L N_1} + E_{N_L N_2} \leq E_{N_L N_1} + E_{N_L N_2}$. If not, the child-flipping algorithm swaps the two children of current node. The meaning of making this swap is : after renumbering the vertexes based on an in-order traversal, if $E_{N_L N_1} + E_{N_L N_2} < E_{N_L N_1} + E_{N_L N_2}$, then

the space to encode the edges $N_L N_2$ and $N_1 N_R$ is greater than that to encode $N_L N_2$ and $N_1 N_R$, because the difference $v_{N_2} - v_{N_L}$ is obviously greater than $v_{N_1} - v_{N_L}$, similar fact works on the difference of $v_{N_R} - v_{N_1}$ and $v_{N_R} - v_{N_2}$.

This heuristic algorithm can be applied to any edge separator tree as a postprocessing step. However, this process is quite time-consuming, the time cost of applying both METIS and child-flipping in Blandfor *etal.* [3]'s experiment shows a demonstration on this. To get the four numbers of edges ($E_{N_L N_1}$, $E_{N_2 N_R}$, $E_{N_L N_2}$ and $E_{N_1 N_R}$), we need to have the four vertex sets: N_L , N_R , N_1 and N_2 for each tree node. The N_L and N_R of each node in separator tree can be pre-collected, but the N_1 and N_2 need to be calculate locally when reach the node. In our project, we make some efforts to reduce the time to conduct this child-flipping processing: Let assume we have a node N_a in a separator tree with left child as N_b and right child as N_c , if we knew the N_1 and N_2 of N_a (denote as N_{1a} and N_{2a}), then

- for node b : $N_{1b} = N_{1a}$, $N_{2b} = N_{2a} \cup N_{Ra}$
- for node c : $N_{1c} = N_{1a} \cup N_{La}$, $N_{2c} = N_{2a}$

Having this general case, then we can generate the N_1 and N_2 for each node during the child-flipping traversal via DFS. The N_L and N_R of each node in separator tree can be pre-collected via a post-order traversal along the edge separator tree, the vertex set of a node is formed by merging the two vertex sets of its children in sorted order. When needing to expand the current N_1 or N_2 , we include the new coming N_L or N_R as distinct sorted set instead of merging N_L into N_1 or N_R into N_2 . Note that

in term of a node, the size of N_L or N_R is always smaller than the size of N_1 or N_2 , respectively. We check an existing of an edge by performing binary on larger sorted set, all of above approaches improve the speed of child-flipping processing significantly

.

Chapter 3

Index Structure

We already introduce how to encode a graph with good edge separator into difference adjacency table bounded by $O(n)$ bits, but to access to the particular adjacency list for a particular vertex, we need to know the exact starting position in the bit sequence. Hence our representation requires a data structure that support efficient select operations on the bit sequence. Note that, the space cost of the data structure should not exceed our space bound which is $O(n)$.

Blandfor *etal.* has proposed three kinds of index structure: *direct*, *semi-direct*, and *indirect* [3]. In addition to these three structures, we also implement other two index structures based on RRR and SD Vector respectively, we will henceforth call these two RRR-based and SDV-based index structure.

direct index structure. The easiest way to access to the starting position of any adjacency list is to store an offset pointer for each vertex. Because our representation uses $O(n)$ bits to store a separable graph, one offset pointer will use $\Theta(\log(n))$ bits. For all vertexes, it will use $n\Theta(\log(n))$ in total. If we want to locate the starting position of any vertex, only one memory access is needed, so this approach is fastest among the five index structures. In our implementation, we allocate one word (32

bits) space for each vertex, hence the total space for this direct index structure is $4n$ bytes.

semi-direct index structure. This *semi-direct* index structure is very similar to the *direct* structure, but instead of allocating one word space for each vertex, we use two words to store four offset values. Let say we have four vertexes $v_n, v_{n+1}, v_{n+2}, v_{n+3}$, The first word space will store the pointer value of v_n , then use the second word to store three offset value $v_{n+1}-v_n, v_{n+2}-v_n, v_{n+3}-v_n$. If the space of the three offset values exceeds one-word bound, then store the offsets somewhere else, then the second word store a pointer pointed to them. Compared to the *direct* structure, if the vertex satisfies $v_{id} \bmod 4 = 0$, then it works exactly as same as *direct* structure, if not, it firstly gets the real offset by adding the corresponding offset to the first pointer value, then access the memory directly. If every four vertexes fit into two-word space, then this *semi-direct* structure would use half of the space used by *direct* structure.

indirect index structure. Note that the above two index structure use $n\theta(n)$, which exceeds $O(n)$ space bound. To limit the space cost within the bound, Bland-*etal.* implement the third structure : *indirect*, which use $O(n)$ space meanwhile supports constant time on locating vertex. The component of the structure is shown in Figure 3.1: we firstly divide all vertexes into blocks with $\log(n)$ vertexes in each, then we further divide each block into subblocks. Each subblock contains a minimal number of vertexes whose indexing offset length reaches at least $k \log(n)$ bits for some constant number k . We store a bit vector with $\log(n)$ bits length for each block, if

v_i is the first vertex in its subblock, then the $(i \bmod \log(n))_{th}$ bit in the vector of its block is set to 1, otherwise 0. We only store the offset pointer of vertex if its corresponding bit in bit vector is 1. This will require $O(n)$ bits in total.

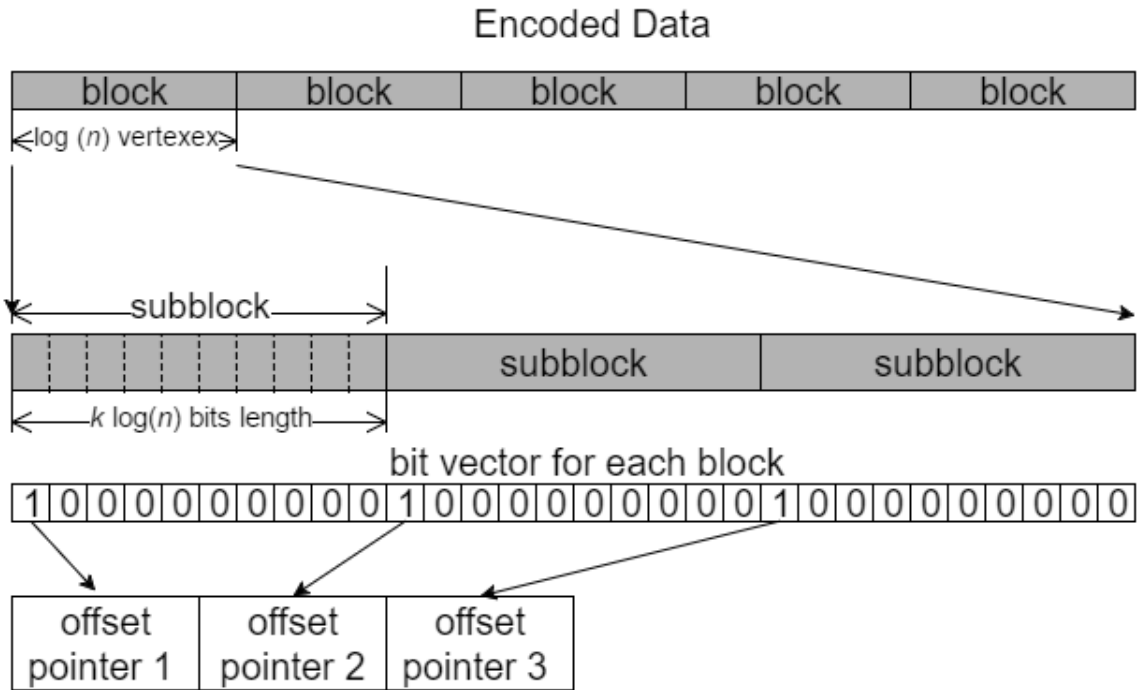


Figure 3.1: The architecture of the *indirect* index structure

Chapter 4

Experiment

Table 4.1: Graph List

Graph	Vertexes	Edges	Max Degree	Source
feocean	143437	409593	6	3D mesh
m14b	214765	1679018	40	3D mesh
febody	45087	163734	28	Scotch Graphs
wave	156317	1059331	44	RIACS Grids
558a	110971	741934	26	Cytoscape
144	144649	1074393	26	Cytoscape
bcsstk30	28924	1007284	218	Matrix: HB
bcsstk31	35588	572914	188	Matrix: HB

Table 4.2: The performance of compact mechanisms.

Graph	METIS		METIS-CF		Degree
	Time	Space	Time	Space	
feocean	55.69	8.68	59.51	8.27	0.86
m14b	129.3	5.21	155.37	5.00	0.51
febody	6.33	4.64	7.2	4.25	0.82
wave	77.4	5.76	86.52	5.57	0.55
558a	39.06	5.45	52.33	5.21	0.54
144	67.09	5.32	84.33	5.13	0.52
bcsstk30	7.89	2.29	12.3	2.24	0.17
bcsstk31	6.61	2.98	9.02	2.89	0.30

Table 4.3: The space cost of various indexing schemes.

Graph	Direct	Semi-D	Indirect	RRR		SD Vector	
	Space	Space	With D	D-free	With D	D-free	With D
feocean	5.6	2.8	0.84	1.96	2.03	1.58	1.59
m14b	2.05	1.02	0.42	0.95	1.01	0.64	0.64
febody	4.4	2.2	0.54	1.22	1.33	1.31	1.31
wave	2.36	1.18	0.46	1.07	1.14	0.72	0.72
558a	2.39	1.19	0.47	1.03	1.10	0.70	0.77
114	2.15	1.07	0.43	0.98	1.05	0.67	0.67
bcsstk30	0.46	0.23	0.26	0.33	0.35	0.16	0.16
bcsstk31	0.93	0.46	0.21	0.51	0.54	0.35	0.35

Table 4.4: The performance of various indexing schemes without degree of vertex encoded.

Graph	Direct		Semi-D		RRR		SD Vector	
	D-query	N-list	D-query	N-list	D0query	N-list	D-query	N-list
feocean	48.32	50.64	49.15	50.25	51.50	51.90	49.72	50.93
m14b	134.75	143.25	138.16	145.80	136.07	141.77	131.44	137.72
febody	8.60	9.94	8.70	10.04	12.03	13.40	9.18	10.55
wave	83.23	89.76	83.43	90.93	87.68	92.82	84.44	89.32
558a	52.59	55.67	52.80	55.93	57.53	60.74	53.20	57.66
144	76.27	79.44	76.92	79.50	79.52	82.52	77.17	80.85
bcsstk30	31.11	35.51	31.44	35.88	33.89	38.38	32.36	36.00
bcsstk31	23.74	27.47	23.85	27.58	29.21	31.93	25.71	28.11

Table 4.5: The performance of various indexing schemes with degree of vertex encoded.

G	Direct		Semi-D		Indirect		RRR		SD Vector	
	D-q	N-l	D-q	N-l	D-q	N-l	D-q	N-l	D-q	N-l
feo*	51.09	54.46	51.33	54.48	56.91	59.37	53.83	57.25	51.73	54.29
m14b	139.61	147.43	141.05	148.15	149.06	153.83	146.05	151.70	142.38	146.14
feb*	10.13	12.54	10.75	12.94	16.81	17.50	13.28	15.67	10.63	13.00
wave	90.97	94.22	91.44	94.55	99.96	101.01	95.35	97.40	91.85	95.31
558a	57.75	62.62	58.13	62.83	66.68	68.65	62.13	66.13	58.97	62.12
114	84.05	87.66	84.24	88.13	91.58	93.64	87.67	81.10	85.56	88.30
b*30	34.08	38.57	34.59	38.60	42.40	45.68	35.49	40.57	34.27	39.81
b*31	26.51	31.02	26.82	31.26	34.78	36.39	30.99	35.78	26.37	31.71

Chapter 5

Conclusion

Did it!

Bibliography

- [1] P. Elias. Universal codeword sets and representations of the integers. *IEEE Transactions on Information Theory*, (21):194–203, 1975.
- [2] Guy Joseph Jacobson. Succinct static data structures. *Carnegie Mellon University, Pittsburgh, PA*, 1988.
- [3] Daniel K.Blandfor, Guy E. Blelloch, and Ian A. Ksh. Compact representations of separable graphs. *Proc. SODA, ACM/SIAM*, page 679–688, 2003.
- [4] R. J. Lipton, D. J. Rose, and R. E. Tarjan. Generalized nested dissection. *SIAM Journal on Numerical Analysis*, (16):346–358, 1979.
- [5] G. L. Miller, S.-H. Teng, W. P. Thurston, and S. A.Vavasis. Separators for sphere-packings and nearest neighbor graphs. *Journal of the ACM*, (44):1–29, 1997.
- [6] D. Okanohara and K. Sadakane. Practical entropy-compressed rank/select dictionary. *Proceedings of ALENEX*, 2007.
- [7] R. Raman, V. Raman, and S. Rao. Succinct indexable dictionaries with applications to representations of k-ary trees and multi-sets. *SODA*, 2002.
- [8] R.J.Liton and R.E. Tarjan. A separator theorem for planar graphs. *SIAM J.Applied Mathematics*, (16):177–189, 1979.